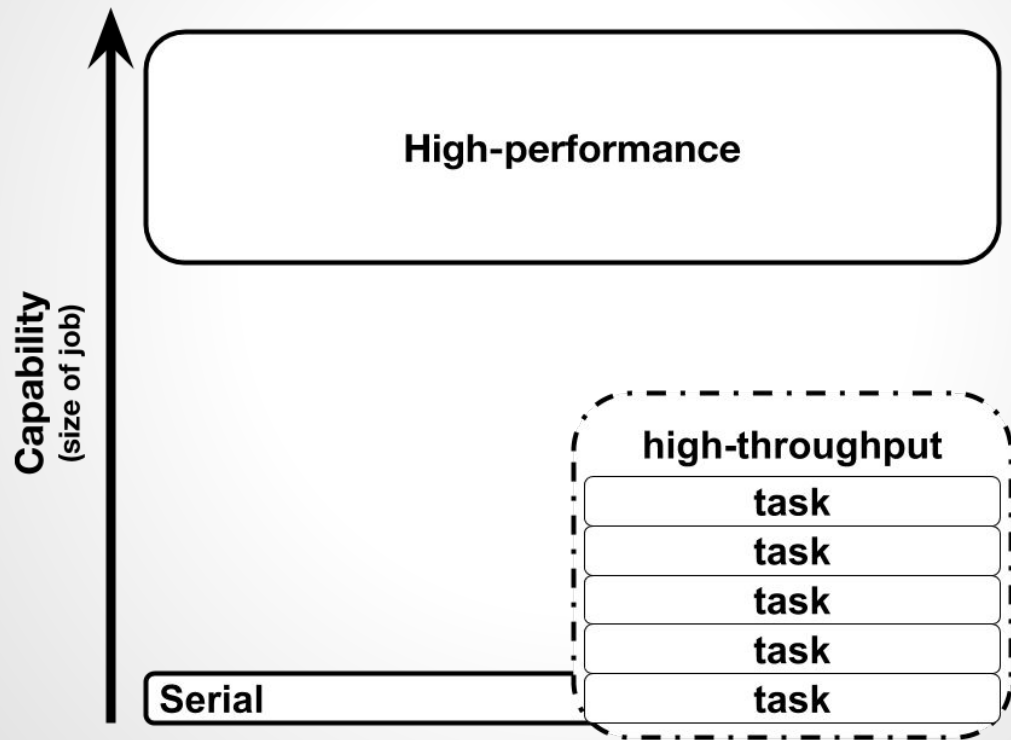


# Efficient Submission of Serial Jobs

Aaron Holt

Slides and Examples: [https://github.com/AaronTHolt/HTC\\_RMACC\\_17](https://github.com/AaronTHolt/HTC_RMACC_17)

# HPC vs HTC

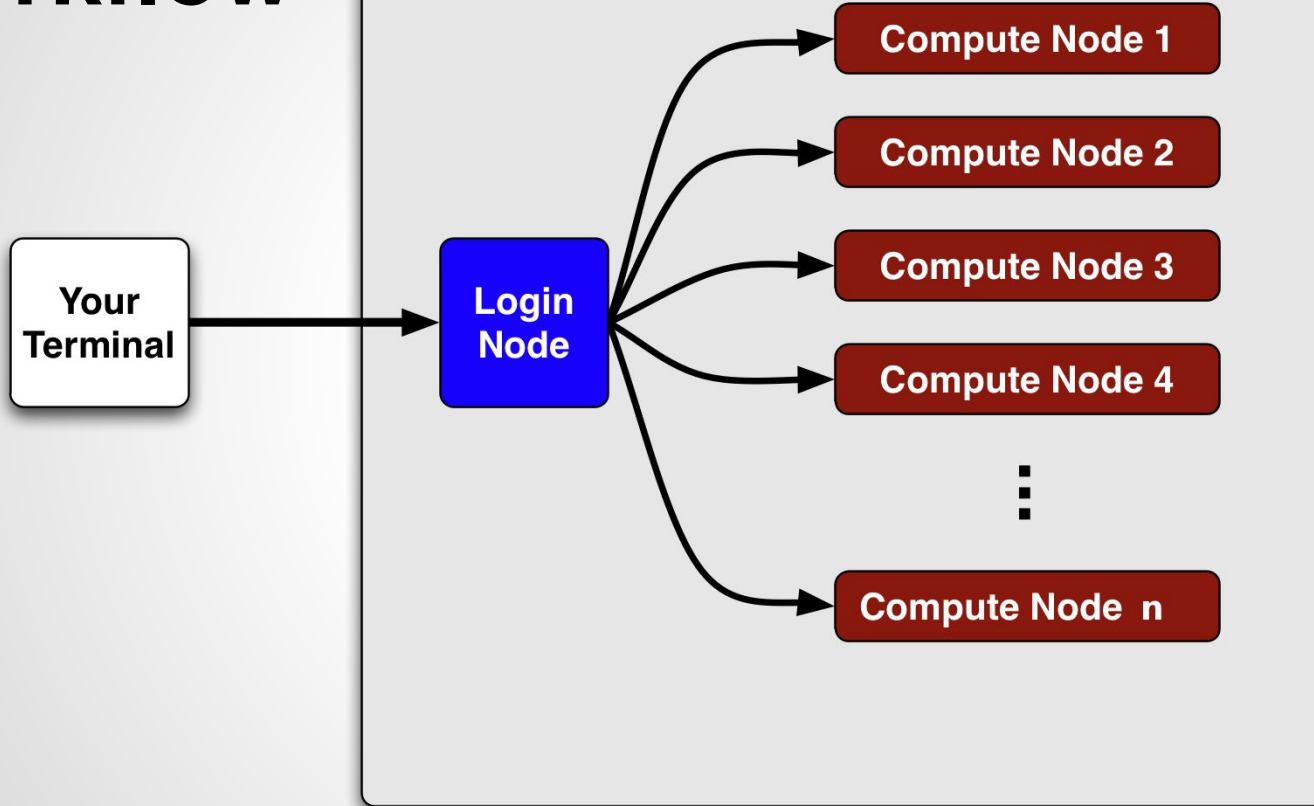


# Tools



CURC loadbalancer

# Workflow



# Setup

- Login to Sandstone HPC (Chrome and Firefox supported)
  - <https://sandstone.rc.colorado.edu/>
  - Start My Server
  - Start
- Those not using Sandstone can login with an ssh client
  - `ssh USERNAME@login.rc.colorado.edu`

# Setup 2

- In the terminal app
  - `git clone https://github.com/AaronTHolt/HTC_RMACC_17`
  - `'cd'` into cloned directory
  - `'module load slurm/summit'`

# Batch Script with One Serial Task

```
#!/bin/bash
#SBATCH --nodes 1
#SBATCH --partition sknl
#SBATCH --output process_file.out
#SBATCH --time 01:30:00

python main.py input_file_1.csv
```

# Batch Script with Multiple Tasks

| Serial Scripts on One Processor   | Serial Scripts on Multiple Processors   |
|---|---|
| <pre>#!/bin/bash #SBATCH --partition sknl #SBATCH --nodes 1 #SBATCH --ntasks-per-node 1 #SBATCH --time 18:00:00  python main.py input_file_1.csv python main.py input_file_2.csv ... python main.py input_file_12.csv</pre> | <pre>#!/bin/bash #SBATCH --partition sknl #SBATCH --nodes 1 #SBATCH --ntasks-per-node 12 #SBATCH --time 01:30:00  python main.py input_file_1.csv &amp; python main.py input_file_2.csv &amp; ... python main.py input_file_12.csv &amp; wait</pre> |



# Your Turn

In the Sandstone editor tab, edit the 'background\_exercise.sh' file (located in HTC\_RMACC\_17/bash). Edit the file to accomplish the following:

1. Commands should run in parallel (use '&' to background)
2. 2 minute wall time
3. Use 1 node in the sknl partition
4. Modify --ntasks-per-node so that each task has a cpu
5. After editing: In the terminal, navigate to the 'bash' directory inside your cloned git directory ('cd bash'). Submit your job with 'sbatch background\_exercise.sh'

# Solution

background\_exercise.sh

```
#!/bin/bash
#SBATCH --nodes 1
#SBATCH --ntasks-per-node 6
#SBATCH --time 0:02:00
#SBATCH --partition sknl
#SBATCH --reservation htc-tutorial
./matrix_mul &
./matrix_mul &
./matrix_mul &
./matrix_mul &
./matrix_mul &
./matrix_mul &
wait
```

Submit with:

sbatch background\_exercise.sh

# Bash Script Summary

- You don't need a special tool
- Available almost everywhere
- Takes some experience to write more complex scripts
- Not great for running large numbers of tasks

# Slurm Job Arrays

- Submit multiple sub-jobs from a single job script
- Array indices specified by `--array`
  - For example, `#SBATCH --array 0-9,12,15`
- Array index is available as `$SLURM_ARRAY_TASK_ID` (in the job)
- Master job id is available as `$SLURM_ARRAY_JOB_ID` (in the job)

# Slurm Job Array 'Hello World'

```
#!/bin/bash
```

```
#SBATCH --nodes 1
```

```
#SBATCH --partition sknl
```

```
#SBATCH --array 0-2
```

```
#SBATCH --output slurm-array-%A.%a.out
```

```
#SBATCH --reservation htc-tutorial
```

```
echo "Master job id: ${SLURM_ARRAY_JOB_ID}"
```

```
echo "Array index: ${SLURM_ARRAY_TASK_ID}"
```

# Submit a Slurm Job Array

- In the terminal app, navigate to  
‘HTC\_RMACC\_17/slurm\_arrays’
- Submit a job with sbatch in the terminal:
  - `sbatch simple_job_array.sh`

# Exercise - Slurm Job Array

There are 10 files called `file_{1..10}.txt` located in `'HTC_RMACC_17/slurm_arrays/data'`. Edit `array_exercise.sh` (in `slurm_arrays` directory) so that the `print_file_contents.py` python script processes each file. For example, the first job index should run the following:

- `python print_file_contents.py data/file_1.txt`

Hints:

- Replace the file number with `${SLURM_ARRAY_TASK_ID}`
- Submit the job with `'sbatch array_exercise.sh'`
- Specify a range of indices with a hyphen (ex: `#SBATCH --array 0-9`)

# Slurm Job Array Summary

- Run on multiple nodes
- No additional tools/software required
- Similar features exist on other schedulers
- Incurs additional scheduling overhead (inefficient for many small tasks)



# Load Balancer

- Submitting hundreds of slurm jobs is inefficient
- RC provides a utility that balances serial applications using MPI (without needing knowledge of MPI!).
- The loadbalancer schedules tasks across multiple nodes after submitting one job
  - Choose how many tasks will run at a time
  - Starts tasks in order (no control over output order)
  - Replaces finished tasks with new ones
  - Straightforward input format

# Load Balancer Create Input File Example

```
for i in {1..100}; do  
    echo "sleep 2; echo process $i" >> cmd_file ;  
done
```

# Load Balancer Input File Example

cmd\_file

sleep 2; echo process 1

sleep 2; echo process 2

sleep 2; echo process 3

...

sleep 2; echo process 98

sleep 2; echo process 99

sleep 2; echo process 100

# Submitting Jobs with lb

submit\_lb.sh

Submitting from the terminal

```
#!/bin/bash
#SBATCH --nodes 3
#SBATCH --ntasks-per-node 6
#SBATCH --time 00:05:00
#SBATCH --output lb-output.out
#SBATCH --partition sknl
#SBATCH --reservation htc-tutorial

module load loadbalance
mpirun lb cmd_file
```

```
$ sbatch submit_lb.sh
```

# Your Turn

Write an input file for the load balancer.

- Input file should be called `cmd_file_2`
- The input file should have at least 2 commands per line
  - One command should be `'hostname'`
  - One command should be `'sleep 2'`
- Example line:
  - `hostname; sleep 2;`
- Should be 50 lines long
- No copy/paste coding! Use a loop.
- Hint: `'echo'` and `'>>'` are useful commands
- Hint: `"` tells bash it's a string and not a command

# Possible Solution

```
for i in {1..50}; do  
    echo "hostname; sleep 2; echo process $i" >> cmd_file_2 ;  
done
```

# Your Turn

Edit the 'exercise\_submit\_lb.sh' batch script and submit a job using the loadbalancer with cmd\_file\_2 as the input file.

- Limit your job runtime to 2 minutes
- Output file loadbalance.out
- 2 nodes
- 3 tasks per node
- Hint: You need to load the 'loadbalance' module
- Hint: 'mpirun lb FILENAME' will runs FILENAME with the loadbalancer

# Possible Solution

exercise\_submit\_lb.sh

```
#!/bin/bash
#SBATCH --nodes 2
#SBATCH --ntasks-per-node 3
#SBATCH --output loadbalance.out
#SBATCH --time 00:02:00
#SBATCH --partition sknl
#SBATCH --reservation htc-tutorial
```

```
module load loadbalance/0.2
mpirun lb cmd_file_2
```

Submit with:

```
sbatch exercise_submit_lb.sh
```



# Output from Multiple Nodes

Ran on 5 nodes with 5 tasks per node.

Input file with “sleep 2; echo process \$i”

process 1

process 2

process 4

process 3

process 5

process 6

process 8

process 7

process 9

# Load Balancer Summary

- No mpi knowledge required
- Saves time by reducing slurm overhead (and queue times for everyone)
- Runs on multiple nodes
- Input file can be created in your favorite language
- Not on other systems (it is on github)

# Summary

- Save yourself some time waiting in the queue by specifying a wall time on your jobs (--time)
- Use the following resources to efficiently submit many serial jobs at once:
  - a) Bash
  - b) Slurm
  - c) CURC Load Balancer

# Questions?

Bash Script: <https://www.rc.colorado.edu/blog/reducejanuswaittimes>

Load Balancer: <https://www.rc.colorado.edu/support/examples-and-tutorials/load-balancer.html>

Slurm: <https://slurm.schedmd.com/>

GNU Parallel

Tutorial: [https://www.gnu.org/software/parallel/parallel\\_tutorial.html](https://www.gnu.org/software/parallel/parallel_tutorial.html)

Examples: <https://www.gnu.org/software/parallel/man.html>

O. Tange (2011): GNU Parallel - The Command-Line Power Tool,  
;login: The USENIX Magazine, February 2011:42-47.

# GNU Parallel

- GNU parallel is a shell tool for executing tasks in parallel using one or more computers.
  - In it's simplest form, GNU parallel is a parallel replacement of a for loop.
- Options to specify how many tasks should run in parallel, display output in order, limit resources and more!
- Multi-node isn't as good as the Load Balancer.

# GNU Parallel Examples

**parallel** [options] [*command* [arguments]] ( ::: arguments | :::+ arguments | ::: argfile(s) | :::+ argfile(s) ) ...

Two ways of printing numbers 1 to 4 in parallel:

```
parallel echo {} ::: 1 2 3 4
```

```
seq 1 4 | parallel echo {}
```

# GNU Parallel Loop Replace Examples

## Bash Loop

```
for i in {1..10}; do  
    echo $i;  
done
```

```
for i in {1..100}; do  
    echo $i | grep 1$;  
done
```

## GNU parallel Replacement

```
seq 1 10 | parallel echo {}
```

```
seq 1 100 | parallel 'echo {} | grep 1$'
```

# Your Turn

Setup:

- On Summit, load GNU parallel
  - `module load gnu_parallel`
- Accept citation agreement
  - `parallel --citation`
  - will cite



# Your Turn

Make this loop run in parallel with GNU parallel:

```
for i in {1..10}; do  
    python print_input.py file_$(i).csv;  
done
```

Hint: use {} instead of \$i:  
file\_{}.csv

```
print_input.py
```

```
import sys  
print(sys.argv)
```

## Previous Example

```
for i in {1..10}; do  
    echo $i;  
done
```

```
seq 1 10 | parallel echo {}
```

# Solution

|              |  |
|--------------|--|
| Original     | <pre>for i in {1..10}; do<br/>    python print_input.py file_\$.csv;<br/>done</pre>  |
| GNU Parallel | <pre>seq 1 10   parallel python print_input.py file_{}.csv</pre>   |
| Output       | <pre>['print_input.py', 'file_1.csv']<br/>['print_input.py', 'file_2.csv']<br/>...<br/>['print_input.py', 'file_10.csv']</pre> |

# GNU Parallel Useful Options

View what commands parallel will run without executing them:

```
$ seq 10 | parallel --dry-run echo {}
```

Limit number of tasks running at one time:

```
$ seq 10 | parallel -j 2 echo {}
```

Wait until enough memory is available to start next task:

```
$ seq 10 | parallel --memfree 2G echo {}
```

See all the options:

```
$ man parallel
```

# GNU Parallel with Slurm

submit\_gnu\_parallel.sh

```
#!/bin/bash
#SBATCH --job-name gnu_parallel
#SBATCH --nodes 1
#SBATCH --output gnu_parallel.out
#SBATCH --time 01:00:00

# The following should be on one line
seq 10 | parallel python print_input.py
file_{}.csv
```

print\_input.py

```
import sys
print(sys.argv) # print command line input

# process data here
```

# GNU Parallel Summary

- Great for replacing loops and speeding them up
  - Control how your tasks are run
  - Can run on multiple computers as well (may take some effort to get working with slurm)
  - Lots of examples and documentation online
  - Useful tool outside of compute nodes too
- 
- Takes time to learn
  - You may have to install a local copy on other systems.

# Additional Problems

Use GNU Parallel to parallelize the following loops:

| Problem 1   | Problem 2  |
|---|--|
| <pre>for color in red green blue ; do     for size in S M L XL XXL ; do         echo \$color \$size     done done</pre> | <pre>(for color in red green blue ; do     for size in S M L XL XXL ; do         echo \$color \$size     done done)   sort</pre> |

# Solutions

|           |   |
|-----------|---|
| Problem 1 | <code>parallel echo {1} {2} ::: red green blue ::: S M L XL XXL</code>        |
| Problem 2 | <code>parallel echo {1} {2} ::: red green blue ::: S M L XL XXL   sort</code> |

# GNU Parallel vs Loop

```
$ seq 1 3 | parallel 'echo {}; echo  
$'
```

```
1          #loop number
```

```
27662      #process id
```

```
2
```

```
27663
```

```
3
```

```
27664
```

```
$ for i in {1..3}; do echo $i; echo  
$$; done
```

```
1
```

```
20614
```

```
2
```

```
20614
```

```
3
```

```
20614
```



